


Data Science

Hands on machine learning
with Scikit-learn

Yueh-Lin Tsai
2019 / 07


About me



- Yueh-Lin Tsai
- Education
 - National Cheng Kung University, M.S., Psychology (2013-2015)
 - National Cheng Kung University, B.S., Psychology (2009-2013)
- Present
 - AI Engineer in Taiwan AI Academy

What is machine learning ?

- Extract relations/patterns from data automatically
- Apply those rules to unseen data



$$f(x) = y$$



Machine learning workflow



- Problem definition
- Data exploration / preprocessing
- Build model
- Model evaluation

Type of machine learning

- Supervised learning
 - Regression
 - Classification
- Unsupervised learning
 - Cluster
 - Dimension reduction
- Reinforcement learning

Scikit - learn

- Package for machine learning in python
- What can sklearn do
 - Preprocess
 - ML models
 - Evaluation metrics



scikit-learn
Machine Learning in Python

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying to which category an object belongs to.

Applications: Spam detection, Image recognition.

Algorithms: SVM, nearest neighbors, random forest, ... — Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.
Algorithms: SVR, ridge regression, Lasso, ... — Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes
Algorithms: k-Means, spectral clustering, mean-shift, ... — Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency
Algorithms: PCA, feature selection, non-negative matrix factorization. — Examples

Model selection

Comparing, validating and choosing parameters and models.

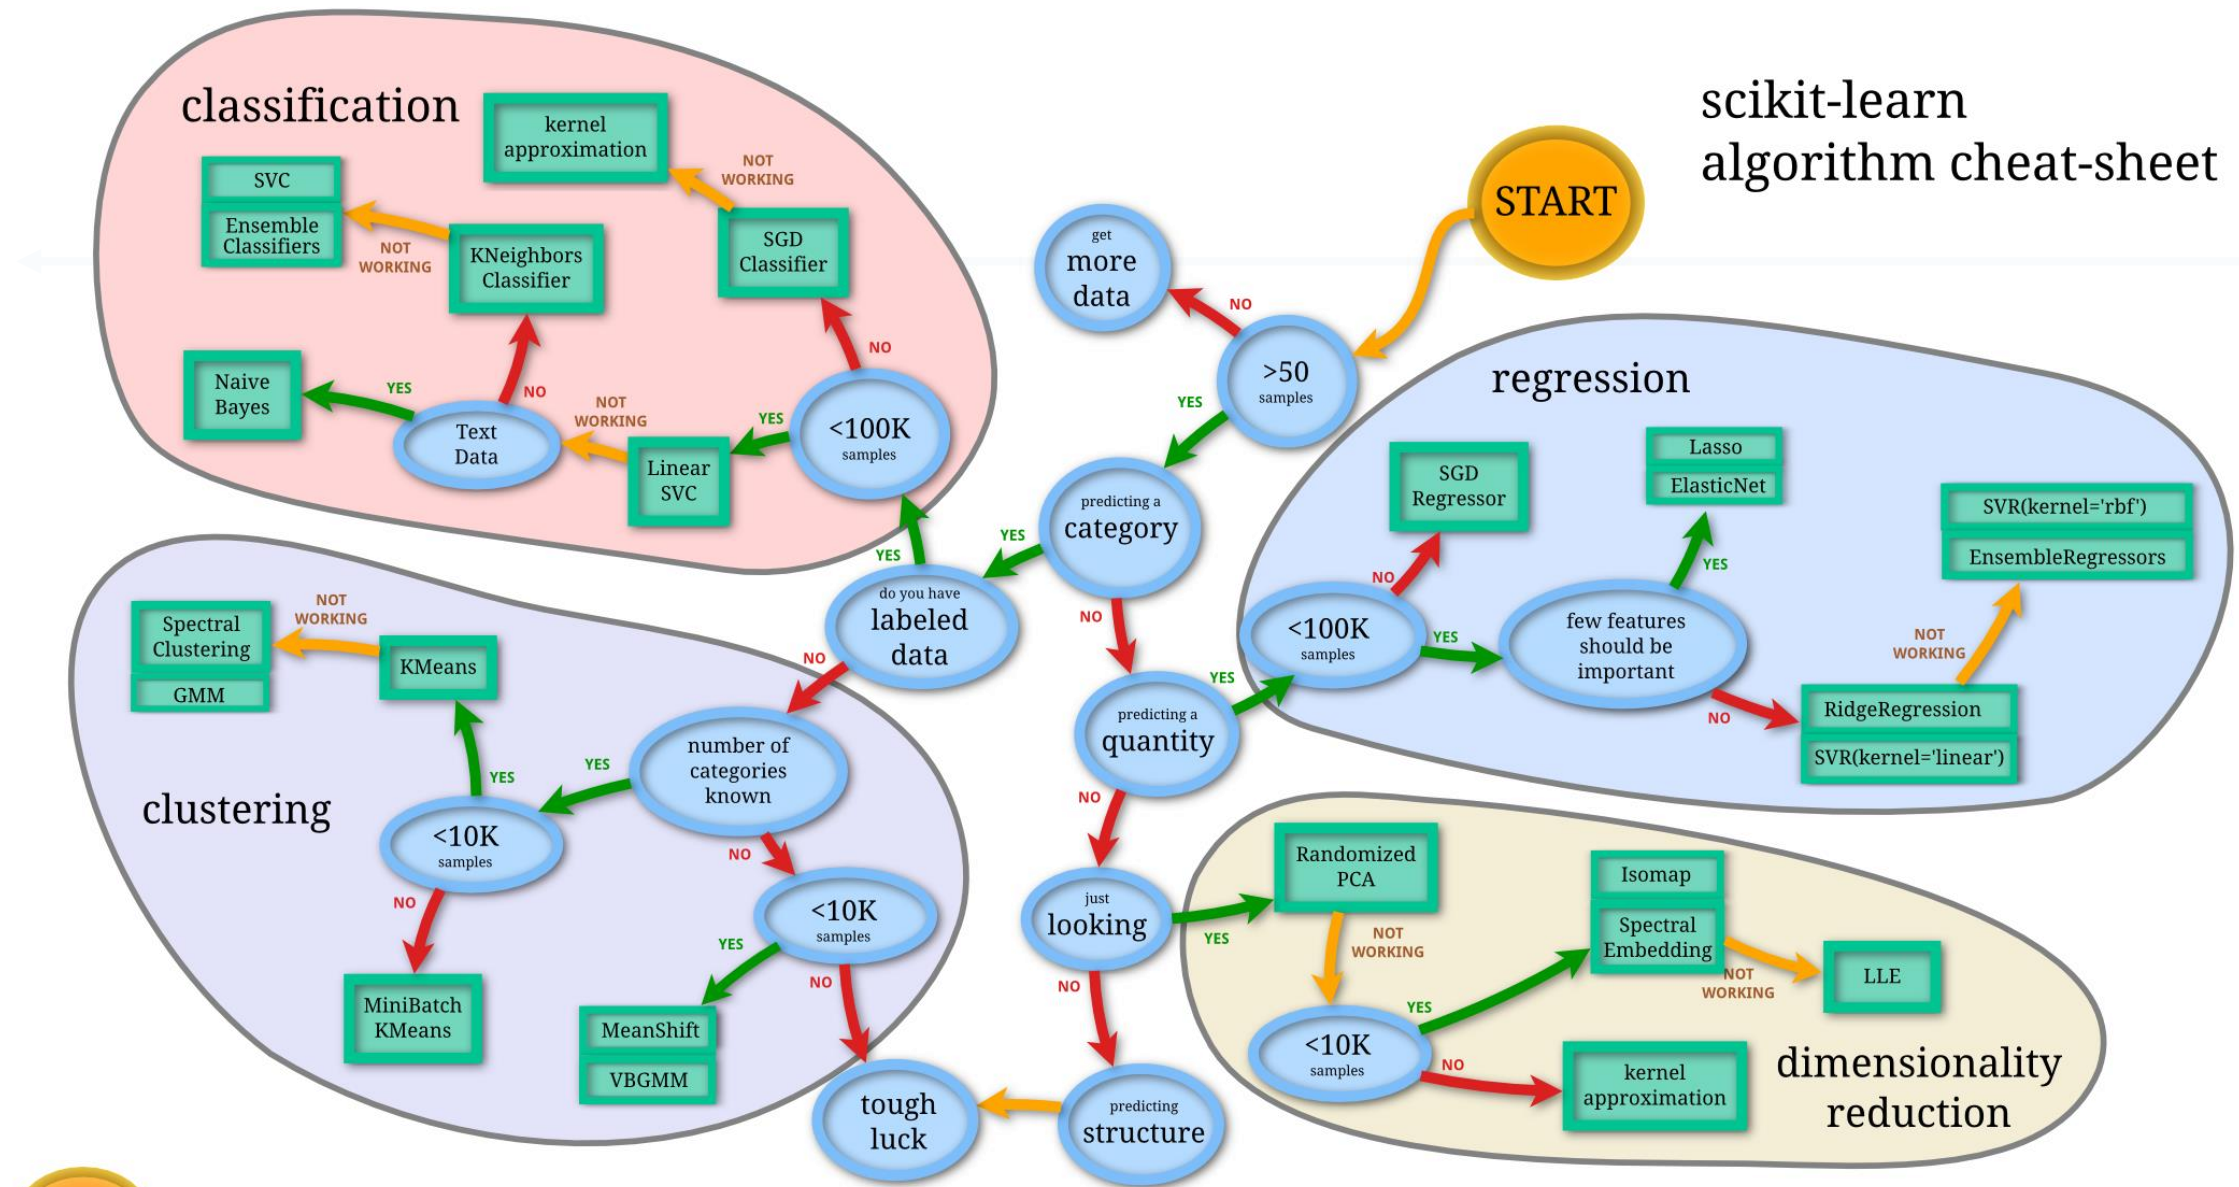
Goal: Improved accuracy via parameter tuning
Modules: grid search, cross validation, metrics. — Examples

Preprocessing

Feature extraction and normalization.

Application: Transforming input data such as text for use with machine learning algorithms.
Modules: preprocessing, feature extraction. — Examples

scikit-learn algorithm cheat-sheet



Build up your ML model in one slide



```
import pandas as pd
from sklearn import preprocessing, linear_model, model_selection, metrics
```

```
data = pd.read_csv('example_data.csv')

data_y = data['target']
data = data.drop('target', axis = 1, inplace = True)
```

```
one_hot_data = pd.get_dummies(data)

ss = preprocessing.StandardScaler()
scale_data = ss.fit_transform(data)
```

```
train_x, test_x, train_y, test_y = model_selection.train_test_split(data, data_y, test_size = 0.2, random_state = 99)
```

```
model = linear_model.LinearRegression() # LogisticRegression()
model.fit(train_x, train_y)

test_prediction = model.predict(test_x)
print('r-square of linear regression : {:.3f}'.format(metrics.r2_score(test_prediction, test_y)))
```


Supervised learning

Machine learning workflow



- Problem definition
- Data exploration / preprocessing
- Build model
- Model evaluation

Data preprocessing



- Handle missing data
 - Delete data which have missing values (row or column)
 - Missing imputation
- Handle outliers
 - Distribution transformation
 - Replace outliers

Data preprocessing

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

- Convert categorical data to numerical data
 - Label encoding
 - One-hot encoding

Name	Score
Amy	78
Bob	90
Chris	65
Amy	86
Chris	67

Name_label	Score
1	78
2	90
3	65
1	86
3	67

Label encoding

Amy_oh	Bob_oh	Chris_oh	Score
1	0	0	78
0	1	0	90
0	0	1	65
1	0	0	86
0	0	1	67

One-hot encoding

Data preprocessing

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

- Normalize data
 - Standard scale
 - Min-max scale

$$x_{standard} = \frac{x - \mu}{\sigma}$$

$$x_{minmax} = \frac{x - Min(x)}{Max(x) - Min(x)}$$

Name	Score
Amy	78
Bob	90
Chris	65
Amy	86
Chris	67

Name	Score
Amy	0.0719
Bob	1.1509
Chris	-1.0969
Amy	0.7912
Chris	-0.9171

Standard scale

Name	Score
Amy	0.48
Bob	0
Chris	1
Amy	0.16
Chris	0.92

Min-max scale

Data preprocessing

```
from sklearn.preprocessing import train_test_split
```

- Data splitting
 - Training set
 - Validation set
 - Testing set
- Cross validation



Model selection



```
from sklearn.linear_model import LinearRegression, LogisticRegression
```


- Linear model
 - Linear Regression
 - Logistic Regression
 - Ridge regression
 - Lasso regression

$$y = f(a_0 + a_1x_1 + a_2x_2 + \dots)$$

The aim of linear model is to find a line which minimize the distance (error) between data and the line

Linear model



- Linear regression



Linear model

$$f(x) = \frac{1}{1 + e^{-(x)}}$$

- Logistic regression

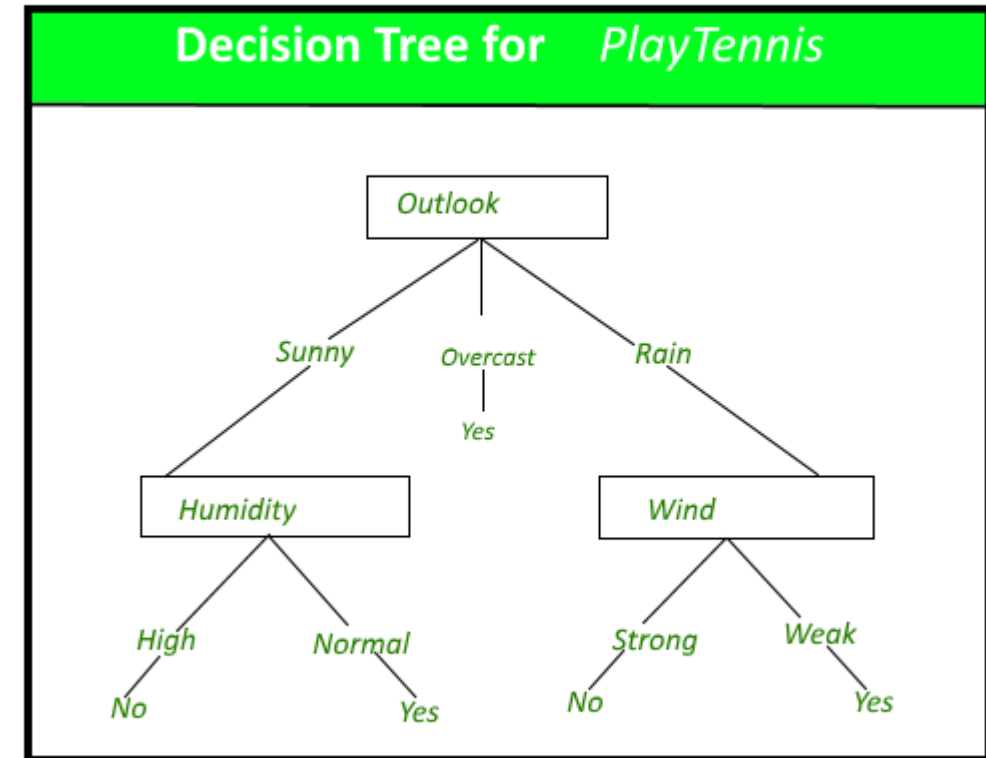


Model selection

```
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor  
from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
```

- Tree-based model
 - Decision Tree
 - Random Forest

The aim of tree-based model is to find good cut-off points to split data repetitively.



Model selection

- Model comparison

- Linear model

- Focus on global information
 - Have data hypothesis

Encoding matters

Normalization is needed

- Tree-based model

- Clear rules provided by model
 - Focus on local information

No need to normalize data

Model selection



- Other machine learning models
 - Support Vector Machine
 - K-Nearest Neighbor
 - Naïve-bayes
 - Neural network

Model evaluation



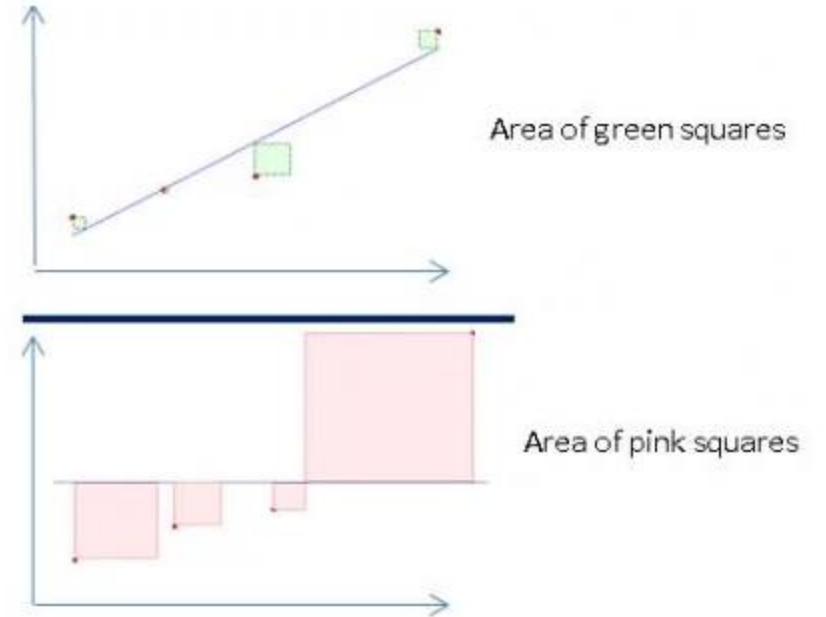
```
from sklearn.metrics import r2_score, mean_squared_error,  
mean_absolute_error
```

- Regression problem
 - R - square, R^2
 - Mean Squared Error, MSE
 - Mean Absolute Error, MAE

Model evaluation

- R-square
 - Range from 0 to 1
 - More is better

$$R^2 = 1 -$$



Model evaluation

- Mean Square Error
- Mean Absolute Error
 - Less is better

$$MSE = \frac{1}{n} \sum \left(\underbrace{y - \hat{y}}_{\substack{\text{The square of the difference} \\ \text{between actual and} \\ \text{predicted}}} \right)^2$$

$$MAE = \frac{1}{n} \sum \left| \underbrace{y - \hat{y}}_{\substack{\text{The absolute value of the} \\ \text{residual}}} \right|$$

Diagram annotations for MAE:

- Divide by the total number of data points**: Points to the $\frac{1}{n}$ term.
- Actual output value**: Points to the y term (green box).
- Predicted output value**: Points to the \hat{y} term (orange box).
- Sum of**: Points to the summation symbol \sum .
- The absolute value of the residual**: Points to the absolute value bars $| \dots |$.

Model evaluation



```
from sklearn.metrics import confusion_matrix, accuracy_score,  
precision_score, recall_score
```


- Classification problem
 - Confusion matrix
 - Accuracy
 - Precision, recall
 - Area under curve (AUC), f1 score

Model evaluation

		PREDICTED	
		Positive	Negative
ACTUAL	Positive	TP True Positive	FN False Negative
	Negative	FP False Positive	TN True Negative

belleva

Type I error
(false positive)



Type II error
(false negative)



Model evaluation

		Actual	
		Positives(1)	Negatives(0)
Predicted	Positives(1)	TP	FP
	Negatives(0)	FN	TN

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + FN + TN}$$

		PREDICTED	
		Positive	Negative
ACTUAL	Positive	True Positive TP	False Negative FN
	Negative	False Positive FP	True Negative TN


Model evaluation

- Precision & recall

$$\text{Precision} = \frac{\text{True Positive}}{\text{Actual Results}} \quad \text{or} \quad \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}}$$

$$\text{Recall} = \frac{\text{True Positive}}{\text{Predicted Results}} \quad \text{or} \quad \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$$

$$\text{Accuracy} = \frac{\text{True Positive} + \text{True Negative}}{\text{Total}}$$



Other topics – supervised learning

- Feature engineering and feature selection
- Overfitting and generalizations
- Parameter selection

Unsupervised learning

Machine learning workflow



- Problem definition
- Data exploration / preprocessing
- Build model
- Model evaluation


Problem definition and model selection

- clustering
 - K-means clustering
 - Hierarchical clustering
 - DBSCAN


The aim of clustering is to categorize data base on their similarity

K-means


from sklearn.cluster import KMeans




0a. Données d'entrée




0b. Initialisation




1a. assignation




1b. calcul des points moyens




2a. assignation




2b. calcul des points moyens



3a. assignation



3b. calcul des points moyens





4a. assignation
clusters stables (fin)

Hierarchical clustering

```
from sklearn.cluster import AgglomerativeClustering
```




階層式



DBSCAN

```
from sklearn.cluster import DBSCAN
```




Problem definition and model selection

- dimension reduction
 - Principal component analysis, PCA

The aim of dimension reduction is to describe high-dimension data with fewer dimension

```
from sklearn.decomposition import PCA
```

PCA



Model evaluation



- Cluster
 - Homogeneity scores
- Dimension reduction
 - Information loss ratio

Supplementary

Feature engineering / feature selection

- Feature generation
 - From domain knowledge
 - By exploring data
- Feature selection
 - Based on correlation
 - Based on coefficients in lasso / ridge regression model
 - Based on feature importance index in tree-based models


Other ML models

- ML models that usually showed on ml competitions
 - Bagging : [Random Forest](#)
 - Boosting : [XGBoost](#), [LightGBM](#), [CatBoost](#)
 - Neural Network
 - Stacking model

Overfitting and generalization

We can always get 100% accuracy on training set with a powerful model, but ...

- How to detect overfitting ?
- How to resolve / avoid overfitting ?
 - Decrease the power of present model
 - Ensemble methods



Parameter selection




```
from sklearn.model_selection import GridSearchCV
```

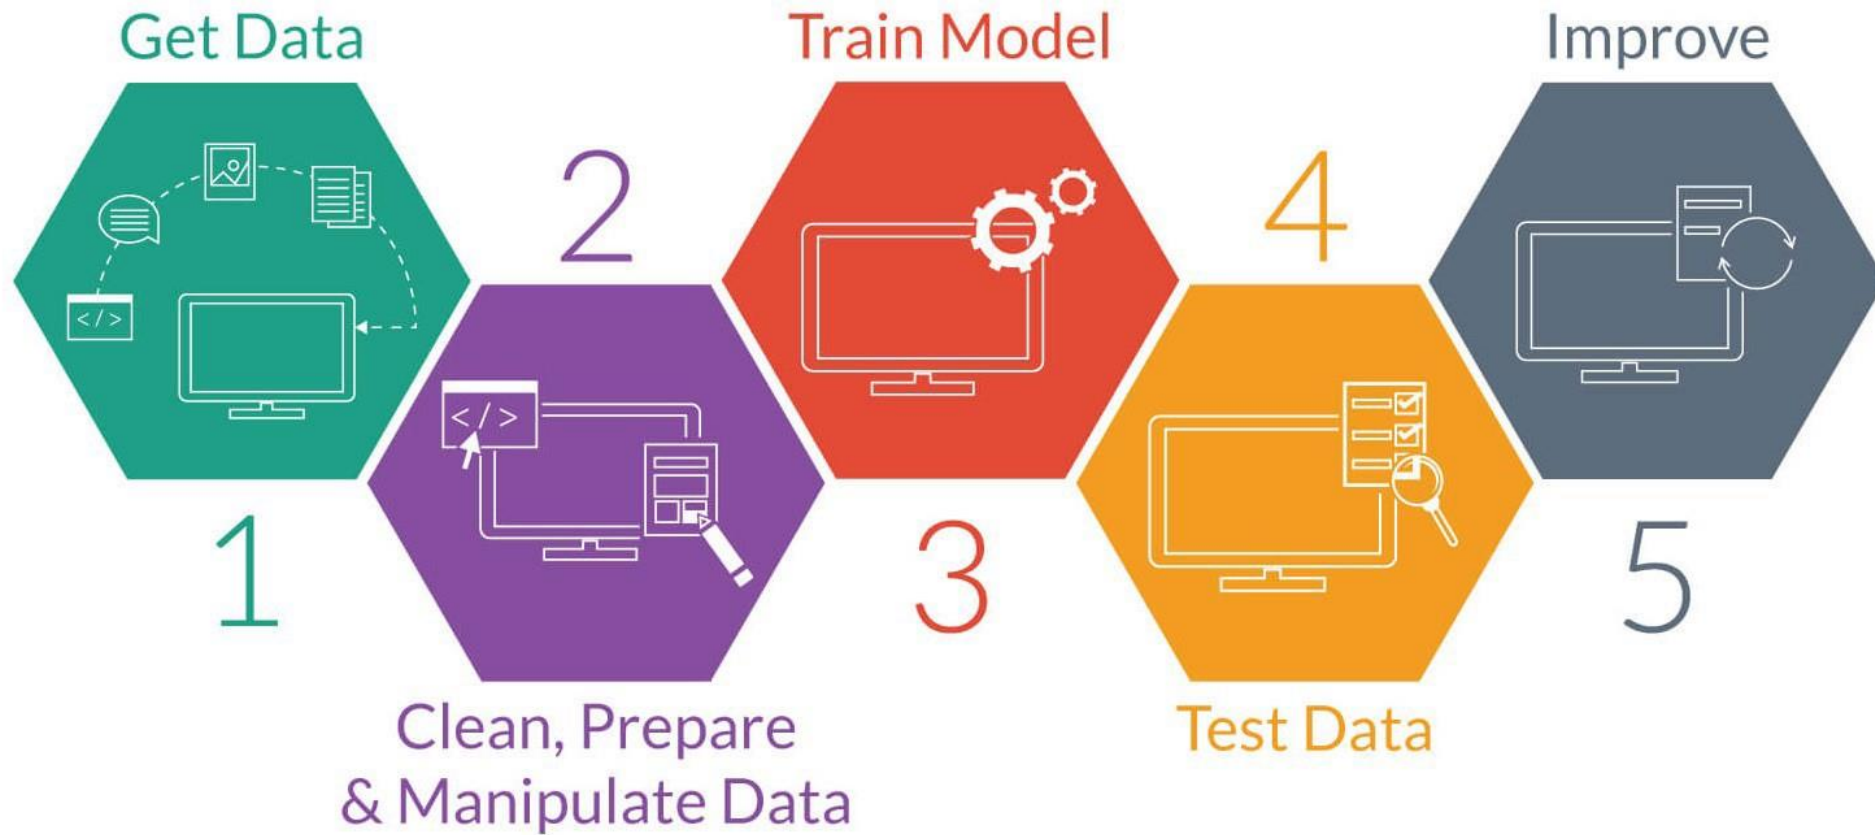
- Select a proper set of parameters would affect the power of model hence make a better prediction
- How to decide which parameter set should be used ?

Parameter selection

Cross Validation



Summary of Machine learning



How to build a great model ?

- Before model fitting

- Preprocessing
- Data exploration
- Data engineering

- After preliminary model

- Parameter tuning
- Cross validation

